

Movie Recommender

Ross Jackson

2023-03-07

```
library(arules)

## Loading required package: Matrix
##
## Attaching package: 'arules'
## The following objects are masked from 'package:base':
##
##      abbreviate, write
movies <- read.transactions("data_movielens_hw1.csv", format = "single",
                           sep = ",", cols = c("userId", "title"), header = TRUE)
```

Obtaining a set of rules of manageable size. In doing so, taking into consideration that the company is mainly interested in developing a system that recommends a movie according to its association with at least two other movies. Plotting the data in order to get a grasp of its characteristics.

```
inspect(movies[1:3])
```

##	items	transactionID
## [1]	{Logan,	
##	The Fundamentals of Caring}	1000
## [2]	{Arrival,	
##	Baby Driver,	
##	Blade Runner 2049,	
##	Call Me by Your Name,	
##	First Reformed,	
##	Get Out,	
##	Hereditary,	
##	Lady Bird,	
##	Mandy,	
##	Mother!,	
##	Phantom Thread,	
##	Piper,	
##	The Handmaiden,	
##	The Neon Demon,	
##	Three Billboards Outside Ebbing, Missouri}	10000
## [3]	{Blade Runner 2049,	
##	Get Out,	
##	Lovesong}	100020

```
summary(movies)


## transactions as itemMatrix in sparse format with
## 24857 rows (elements/itemsets/transactions) and
```

```
## 3222 columns (items) and a density of 0.002046879
##
## most frequent items:
##           Deadpool           Arrival
##           8322           6055
##           Zootopia Rogue One: A Star Wars Story
##           5341           3374
##           Logan           (Other)
##           3344           137497
##
## element (itemset/transaction) length distribution:
## sizes
##      1      2      3      4      5      6      7      8      9     10     11     12     13     14     15     16
## 7648 3962 2460 1807 1249  916  843  643  568  442  395  386  308  281  245  236
##    17    18    19    20    21    22    23    24    25    26    27    28    29    30    31    32
##   197   183   158   126   152   113   102   101    95    85    82    71    65    59    62    52
##    33    34    35    36    37    38    39    40    41    42    43    44    45    46    47    48
##    47    45    35    37    39    27    37    25    27    33    17    27    28    15    17    11
##    49    50    51    52    53    54    55    56    57    58    59    60    61    62    63    64
##    22    15     8     7    15    13    13     9    16    12     9     8     8     9     5     3
##    65    66    67    68    69    70    71    72    73    74    75    76    77    78    79    80
##     3     9     9     4     3     4     5     7     4     2     3     1     4     4     6     3
##    81    82    83    84    85    86    87    88    90    91    93    94    96    97   100   103
##     1     3     1     5     2     5     3     1     3     2     3     2     1     2     2     1
##   104   106   108   109   111   112   113   117   119   122   123   138   161   182   183   193
##     1     1     1     1     1     3     1     1     1     1     1     1     1     1     1     1
##
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
##    1.000   1.000   3.000   6.595   7.000 193.000
##
## includes extended item information - examples:
##      labels
## 1  '63 Boycott
## 2  #realityhigh
## 3 1 Mile to You
##
## includes extended transaction information - examples:
##      transactionID
## 1           1000
## 2          10000
## 3          10020
```

```
movie_freq <- itemFrequency(movies)
head(movie_freq)
```

```
##           '63 Boycott           #realityhigh           1 Mile to You           1 Night
##           4.023012e-05           2.816108e-04           4.023012e-05           1.609205e-04
##                1:54 10 Cloverfield Lane
##           8.046023e-05           6.521302e-02
```

```
itemFrequencyPlot(movies, topN = 5)
```



Parameter choice explanation: support: too many rules returned when support= 0.01, had to be increased
 confidence: Will return rules with higher association with eachother minlen = 2, as it is specified that
 “take into consideration that the company is mainly interested in developing a system that recommends a
 movie according to its association with at least two other movies.” maxlen = 8 as warning message returned
 otherwise

```
movie_rules <- apriori(movies, parameter = list(support = 0.02, confidence = 0.8, minlen = 2, maxlen = 8))
```

```
## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##           0.8   0.1   1 none FALSE               TRUE     5   0.02     2
## maxlen target ext
##           8 rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##     0.1 TRUE TRUE  FALSE TRUE     2     TRUE
##
## Absolute minimum support count: 497
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[3222 item(s), 24857 transaction(s)] done [0.06s].
## sorting and recoding items ... [71 item(s)] done [0.00s].
## creating transaction tree ... done [0.01s].
## checking subsets of size 1 2 3 4 5 done [0.01s].
## writing ... [12 rule(s)] done [0.00s].
## creating S4 object ... done [0.01s].
```

```
summary(movie_rules)
```

```
## set of 12 rules
##
```

```

## rule length distribution (lhs + rhs):sizes
## 4
## 12
##
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
##         4         4         4         4         4         4
##
## summary of quality measures:
##      support      confidence      coverage      lift
##  Min.   :0.02003  Min.   :0.8002  Min.   :0.02450  Min.   :2.390
## 1st Qu.:0.02035  1st Qu.:0.8054  1st Qu.:0.02512  1st Qu.:2.406
##  Median :0.02124  Median :0.8068  Median :0.02639  Median :2.459
##   Mean   :0.02216  Mean   :0.8133  Mean   :0.02724  Mean   :3.815
## 3rd Qu.:0.02362  3rd Qu.:0.8203  3rd Qu.:0.02879  3rd Qu.:6.519
##   Max.   :0.02639  Max.   :0.8325  Max.   :0.03243  Max.   :6.701
##      count
##  Min.   :498.0
## 1st Qu.:505.8
##  Median :528.0
##   Mean   :550.8
## 3rd Qu.:587.0
##   Max.   :656.0
##
## mining info:
##      data ntransactions support confidence
##  movies      24857      0.02      0.8
##
##      call
## apriori(data = movies, parameter = list(support = 0.02, confidence = 0.8, minlen = 2, maxlen = 8))
rules1 <- sort(movie_rules, by = "support") # sort by support
inspect(rules1[1:5]) #inspect the top 5 rules

##      lhs      rhs      support confidence      coverage
## [1] {Doctor Strange,
##      Guardians of the Galaxy 2,
##      Logan}      => {Deadpool}      0.02639096  0.8324873 0.03170133 2.
## [2] {Captain America: Civil War,
##      Doctor Strange,
##      Guardians of the Galaxy 2}      => {Deadpool}      0.02594842  0.8002481 0.03242547 2.
## [3] {Captain America: Civil War,
##      Doctor Strange,
##      Logan}      => {Deadpool}      0.02458060  0.8301630 0.02960937 2.
## [4] {Captain America: Civil War,
##      Guardians of the Galaxy 2,
##      Logan}      => {Deadpool}      0.02329324  0.8166432 0.02852315 2.
## [5] {Deadpool,
##      Doctor Strange,
##      Thor: Ragnarok}      => {Guardians of the Galaxy 2} 0.02212656  0.8064516 0.02743694 6.

```

Below I have computed the standardized lifts and bounds for the rules. This allows us to compare the rules on a level playing field.

```

## Standardized Lift:
rules1 <- apriori(movies, parameter = list(support = 0.02, confidence = 0.8, minlen = 2, maxlen = 8))

## Apriori

```

```

##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##      0.8      0.1      1 none FALSE              TRUE      5      0.02      2
## maxlen target  ext
##      8 rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##      0.1 TRUE TRUE  FALSE TRUE      2      TRUE
##
## Absolute minimum support count: 497
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[3222 item(s), 24857 transaction(s)] done [0.05s].
## sorting and recoding items ... [71 item(s)] done [0.00s].
## creating transaction tree ... done [0.01s].
## checking subsets of size 1 2 3 4 5 done [0.01s].
## writing ... [12 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].

qual <- quality(rules1) # extract quality measures

# compute p(A) and p(B)
pA <- qual$coverage
pB <- qual$confidence/qual$lift
# compute lift upper and lower bounds
U <- apply(cbind(1/pA, 1/pB), 1, min)
L <- apply(cbind(1/pA + 1/pB - 1/(pA*pB), 0.01/(pA*pB), 0.5/pB, 0), 1, max)
std_lift <- (qual$lift - L)/(U - L) # standardized lift
data.frame(rule = labels(rules1),
           lift = qual$lift, L, U, std_lift) # print rules and associated metrics

##
## 1 {Deadpool,Thor: Ragnarok,Untitled Spider-Man Reboot} => {Guardians of the Galaxy 2}
## 2 {Doctor Strange,Guardians of the Galaxy 2,Untitled Spider-Man Reboot} => {Deadpool}
## 3 {Deadpool,Doctor Strange,Untitled Spider-Man Reboot} => {Guardians of the Galaxy 2}
## 4 {Captain America: Civil War,Doctor Strange,Thor: Ragnarok} => {Guardians of the Galaxy 2}
## 5 {Deadpool,Doctor Strange,Thor: Ragnarok} => {Guardians of the Galaxy 2}
## 6 {Captain America: Civil War,Doctor Strange,Guardians of the Galaxy 2} => {Deadpool}
## 7 {Captain America: Civil War,Doctor Strange,Logan} => {Deadpool}
## 8 {Captain America: Civil War,Doctor Strange,Zootopia} => {Deadpool}
## 9 {Captain America: Civil War,Logan,Rogue One: A Star Wars Story} => {Deadpool}
## 10 {Captain America: Civil War,Guardians of the Galaxy 2,Logan} => {Deadpool}
## 11 {Doctor Strange,Logan,Rogue One: A Star Wars Story} => {Deadpool}
## 12 {Doctor Strange,Guardians of the Galaxy 2,Logan} => {Deadpool}
## lift L U std_lift
## 1 6.517054 4.045736 8.091471 0.6108453
## 2 2.406644 1.493451 2.986902 0.6114650
## 3 6.700957 4.045736 8.091471 0.6563011
## 4 6.616671 4.045736 8.091471 0.6354680
## 5 6.525380 4.045736 8.091471 0.6129032
## 6 2.390263 1.493451 2.986902 0.6004963
## 7 2.479616 1.493451 2.986902 0.6603261
## 8 2.405340 1.493451 2.986902 0.6105919

```

```
## 9  2.410822 1.493451 2.986902 0.6142626
## 10 2.439233 1.493451 2.986902 0.6332863
## 11 2.402896 1.493451 2.986902 0.6089552
## 12 2.486558 1.493451 2.986902 0.6649746
```

As we can see from this, rules 12, 7 and 3 are 3 of the most interesting, as there standardized lifts are the highest obtained from the set of 12 rules.